Planner-to-Executor Contract Design and Runtime Validation for Embodied-Agent Execution: A Controlled Simulator Study

Anonymous Authors

***Abstract*—We present a controlled study that isolates planner-to-executor contract design and lightweight runtime validation for embodied-agent execution. The main factorial comparisons use a lightweight deterministic reference backend; a minimal Isaac Sim navigation slice provides a limited qualitative cross-backend check that the same ordering appears under physics-backed execution. Across navigation and tabletop manipulation, we compare an under-specified direct-action contract (P0), a typed tool-call contract (P1), and the same typed contract with a brief self-check (P2), under bare dispatch (R0) or validation with a single retry (R1). In the fixed-runtime action-interface ablation, P0/R0 fails on all four evaluated tasks, whereas P1/R0 and P2/R0 succeed on all four and eliminate invalid actions. In the fixed-contract runtime-only ablation, validation plus one retry raises success from 2/4 to 4/4 by converting invalid-first-action runs into successful recoveries. Across the broader fixed task matrix, structured contracts remain reliable in both task families, and P2 is associated with fewer planner/tool calls than P1 in the evaluated comparisons. Because the planner is deterministic, the results are descriptive comparisons over fixed task instances rather than stochastic rollout estimates. These results identify explicit action interfaces and lightweight runtime validation as practical architectural design choices within the tested setting and evaluated task matrix.**

***Index Terms*—embodied agents, robot execution, action interfaces, runtime validation, task planning**

I. Introduction

Embodied-agent systems often fail at the boundary between planning and execution. A planner can identify the right task-level intention yet still emit an action name or argument structure that the executor cannot dispatch. Such failures expose not only planning errors but also the planner-to-executor contract, the action interface that mediates the handoff, and the runtime validation policy that decides whether a bad first emission should terminate the run or be repaired.

This execution boundary often remains implicit in embodied-agent studies, where broader evaluations and stronger end-to-end models dominate the narrative [1]–[5]. Those studies are valuable, but they entangle execution-stack design with backend scale and benchmark breadth. We therefore ask a narrower robotics-systems question: when the simulator setting and task families are held fixed, how much do planner-to-executor contract design and lightweight runtime validation affect invalid actions, recovery, success, and planner/tool call counts?

We study two task families—language-conditioned navigation and tabletop manipulation—using a lightweight deterministic reference backend for the main factorial comparisons and a minimal Isaac Sim navigation slice for qualitative cross-backend checking. The matrix compares three contract variants: an under-specified direct-action contract (P0), a typed tool-call contract (P1), and the same typed contract with a brief self-check (P2). It also compares two runtime variants: bare dispatch (R0) and validation with a single retry (R1). These variants are implemented through different planner instructions, but the manipulated system variables are the action interface presented to the executor and the runtime validation policy applied before dispatch. The planner backend is deterministic by design, so each condition is compared once per task instance rather than through repeated stochastic rollouts.

The study is designed to separate two effects: whether a declared tool interface prevents invalid dispatches at the source, and whether one validation-and-retry step recovers a narrow class of violations that still occur. We include harder tasks to test whether added burden changes the qualitative ordering or mainly increases planner/tool activity. The paper is limited to this controlled simulator study and draws bounded execution-design lessons rather than new-model or benchmark claims.

This paper makes three contributions:

1) It presents a controlled simulator boundary study of planner-to-executor contract design and lightweight runtime validation across navigation and tabletop manipulation. Quantitative comparisons come from a lightweight reference backend, with a small Isaac Sim navigation slice for qualitative cross-backend checking.
2) It provides evidence within the evaluated task matrix that typed tool-call contracts reduce invalid actions relative to an under-specified direct-action contract, while runtime validation with a single retry converts recoverable invalid-first actions into successful executions in both covered task families.
3) It extracts a compact architectural lesson from the tested setup: contract design, runtime validation, and a brief planner-side self-check are separate architectural choices that shape reliability and planner/tool call counts without changing the underlying task matrix.

Fig. 1 summarizes the evaluated boundary: a task instruction reaches the planner, the planner emits an action request under one of the three contract designs, the runtime applies either bare dispatch or validate-and-retry, and the executor records the outcome metrics used throughout the paper. The same

figure also carries the compact failure-repair trace that anchors the recovery discussion.

Table I juxtaposes one planner emission per contract: P0 may still step outside the declared callable namespace, P1 binds the request to a dispatchable typed call, and P2 adds only a brief planner-side self-check.

II. Related Work

A. Embodied and LLM-Based Robot Planning

Recent embodied and language-conditioned robot-planning work studies how high-level language intent is translated into executable robot behavior in simulation or on robots, including closed-loop language feedback, embodied multimodal models, vision-language-action policies, model-based language grounding, and systems-first integration pipelines such as Inner Monologue [1], [2], [6]–[8]. This also includes language-conditioned skill selection, code-generating control, and program-like robot planning directions exemplified by SayCan, Code as Policies, and ProgPrompt [9]–[11]. That literature makes planner quality visible, but it often leaves less room to isolate what happens at the boundary between planner output and dispatch. Our paper is narrower: it holds the simulator setting fixed and studies how planner-to-executor contract design changes invalid actions, recovery, success, and planner/tool call counts.

The comparison to SayCan, Code as Policies, and ProgPrompt is the planner/executor interface itself. SayCan ranks candidate skills through affordance grounding, whereas our manipulated variable is whether the planner must emit one typed tool call from a declared task-scoped interface. Code as Policies synthesizes executable code over robot APIs, while our planner emits one structured call rather than a generated control program. ProgPrompt is closer in its use of programmatic primitives, but our study focuses on the declared per-step action interface and the runtime validation boundary rather than on program generation itself.

B. Execution Contracts, Runtime Validation, and Robotics Middleware

A second line of work addresses execution architecture rather than only planner quality. Relevant threads include ROS 2 middleware and action semantics [12], [13], planning frameworks such as PlanSys2 [14], behavior-tree execution stacks and libraries such as Behavior Trees in Robotics and AI and BehaviorTree.CPP [15], [16], and task-and-motion planning interfaces such as PDDLStream and the broader TAMP literature [17], [18]. The shared intuition is that planner outputs should not be treated as executable by default; they need an explicit interface to the execution stack. The present study is closest to that execution-boundary view. Its runtime layer is intentionally lightweight: R1 validates an emitted action or tool call against the active contract and allows one repair attempt. The contribution is therefore evidence about a small validation-and-retry mechanism in a controlled simulator setting, not a claim of a complete middleware stack or end-to-end safety case.

C. Controlled Design Studies and Broader Evaluations

Methodologically, the paper favors a small controlled matrix over benchmark breadth. Broader embodied-agent evaluations such as EmbodiedBench, IS-Bench, and the Isaac Sim-based Mind and Motion Aligned benchmark are valuable because they expose long-horizon execution difficulty, safety failures, or integrated planning-and-control behavior across larger task sets or evaluation matrices [3]–[5]. By varying only the action interface and runtime validation policy within a fixed simulator setting, we instead use a narrower study design to attribute differences in invalid actions, recovery, success, and planner/tool call counts to explicit execution-design choices. This complements broader embodied-agent evaluations rather than competing with them.

III. Study Design and Experimental Setup

A. Controlled Simulator Setting

The study uses two execution backends. The *primary* backend is a lightweight deterministic reference environment implemented in pure Python; it carries the full factorial evaluation matrix and provides the quantitative comparisons reported in the paper. A *secondary* backend runs a minimal Isaac Sim stage [19] on two navigation instances as a limited qualitative cross-backend check that the same condition ordering (P0 worst, P1/P2 succeed, R1 recovers) also appears under physics-backed execution. The navigation family contains language-conditioned mobile-goal tasks in indoor scenes. The manipulation family contains tabletop tasks that require object identification, grasping, transfer, and placement. The main comparison uses a shared easy slice plus one harder slice per family, letting us test whether additional task burden changes the qualitative ordering or mainly amplifies planner/tool call counts.

All conditions share the same high-level handoff:

1) The planner receives the current task state and the exposed tool set.
2) The planner emits an action request.
3) Under R0, the request is dispatched as issued; under R1, the runtime validates it against the active contract and allows one repair attempt on failure.
4) The executor runs only requests that satisfy the active dispatch interface.

The dispatch vocabulary is family-specific and explicit. Navigation exposes state getters and `navigate_to`; manipulation exposes state getters and `scripted_pick_place_step`. The study therefore asks whether the planner can hand the executor a dispatchable request under a specified interface, not whether it can merely describe the task in natural language.

B. Contract and Runtime Variants

The controlled matrix is summarized in Table II. P0 uses an under-specified direct-action contract, allowing the planner to emit a task-level next action without binding it to the executor's tool namespace. P1 uses a typed tool-call contract,

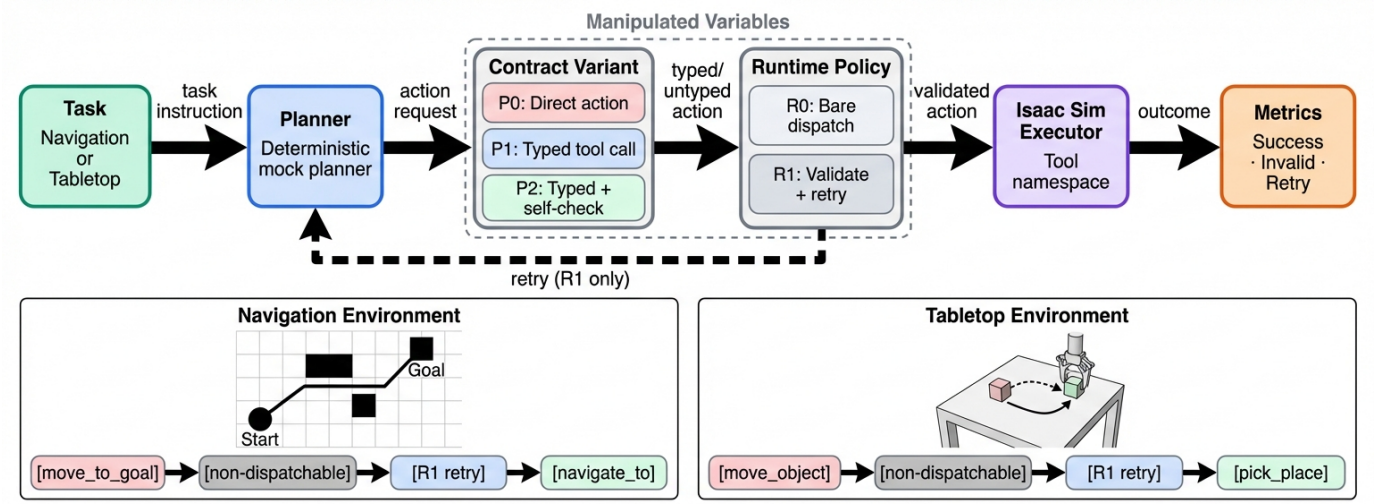

Fig. 1. System overview of the evaluated planner-to-executor boundary. A task is mapped by the planner into one of three contract variants, then passed through bare dispatch `R0` or validate-and-retry `R1` before the executor (a lightweight deterministic backend for the main factorial, with a minimal Isaac Sim navigation slice for qualitative cross-backend checking). Insets show representative navigation and tabletop task scenes together with compact failure-repair traces in which non-dispatchable verbs are rejected and corrected into executor-visible tools.

TABLE I
SCHEMATIC PLANNER OUTPUTS UNDER A SHARED DECLARED TOOL LIST, CONDENSED FROM RECORDED LOGS. P0 MAY EMIT AN UNDECLARED ACTION, P1 YIELDS A DISPATCHABLE TYPED CALL, AND P2 ADDS A BRIEF PLANNER-SIDE SELF-CHECK.

Declared tools: `get_robot_state | get_goal_state | navigate_to(target_pose)`

	P0 — Free JSON minimal action	**P1 — Typed call** one declared tool + args	**P2 — Typed call + self-check** same call + brief self-check

Example emission	`tool_name: move_to_goal`
	`arguments: {}`
		`tool_name: navigate_to`
		`target_pose: (1.0, 0.0, 0.0)`
			`tool_name: navigate_to`
			`target_pose: (1.0, 0.0, 0.0)`
			`self_check: target_pose matches goal pose`

Dispatchability	undeclared tool possible; not dispatchable	declared callable + required args; dispatchable	same call surface as P1; + planner-side check

requiring the planner to choose one exposed tool and provide dispatchable arguments. P2 keeps the same typed tool-call contract and adds a brief planner-side self-check before emission. Table I juxtaposes one planner emission for each contract format.

R0 is a bare-dispatch runtime: it does not perform pre-dispatch validation or offer a repair path. R1 is a validate-and-retry runtime: it checks the emitted request against the active contract and, on failure, returns the validation error to the planner and allows one retry. Contract and runtime therefore play distinct roles. Contract design can prevent invalid actions at the source, while runtime validation can recover some contract violations that still slip through.

C. Implementation Snapshot

The study uses a repo-local deterministic planner backend implemented in `src/isaacsim_agent/planner/mock.py` rather than a hosted foundation-model API. Given the current task state, exposed tools, and any retry message, it emits one JSON action per planner call; there are therefore no temperature, top-p, or rollout-sampling settings in this evaluation. Under P0, the planner emits task-level verbs (`move_to_goal`, `move_object`) that are semantically plausible but absent from the executor's tool registry—*injected boundary probes* designed to expose the contract gap, not stochastic foundation-model errors.

Crucially, the three contract conditions separate two distinct design axes. P0 versus P1 isolates *action-interface specification* under a fixed action-selection policy: both conditions share the same planner decision logic, task-state input, tool registry, and planning strategy; only the serialization format and dispatch interface change. P2 adds one bounded planner-side self-check on top of the same typed interface used by P1, so the P1-versus-P2 comparison isolates a *planner-side deliberation step* rather than a second interface change. Planner instructions for all three contracts are included in the supplementary material and public repository so that readers can verify this separation.

P0 realizes a direct-action JSON contract, P1 realizes typed tool JSON, and P2 keeps the same typed tool JSON while adding a planner-side `self_check` string. The planner request already exposes the valid tool names, descriptions, and

TABLE II
FIXED EVALUATION MATRIX USED IN THE MANUSCRIPT. THE MAIN COMPARISON COVERS 21 TASK INSTANCES ACROSS FOUR COHORTS, EACH EVALUATED UNDER ALL SIX CONTRACT/RUNTIME CELLS; THE TWO COMPACT ABLATIONS ADD FOUR TASK INSTANCES EACH TO ISOLATE CONTRACT AND RUNTIME EFFECTS.

Comparison	Family	Slice	Varying axes	Tasks (= runs/cell)	Total runs	Purpose
Main matrix	Navigation	Easy	P0/P1/P2 x R0/R1	11	66	Main ordering
Main matrix	Manipulation	Easy	P0/P1/P2 x R0/R1	3	18	Main ordering
Main matrix	Navigation	Harder	P0/P1/P2 x R0/R1	4	24	Harder robustness
Main matrix	Manipulation	Harder	P0/P1/P2 x R0/R1	3	18	Harder robustness
Action-interface ablation	Mixed	Easy shared subset	P0/P1/P2 at fixed R0	4	12	Contract mechanism
Runtime-only ablation	Mixed	Recovery probes	R0/R1 at fixed P1	4	8	Runtime recovery

required arguments through `available_tools`. R1 adds one more element: if parsing or validation fails, it appends the literal error string to a short repair instruction and allows one retry. The runtime parses the emitted JSON, validates the tool name against the task-specific registry, enforces the exact required argument set, and then checks task-specific payload semantics. For navigation, `navigate_to` must carry a valid `target_pose` that matches the configured goal pose. For manipulation, `scripted_pick_place_step` must carry the next scripted `phase_name`. The added P2 `self_check` field is recorded in the planner trace but does not replace runtime contract checking.

Navigation exposes `get_robot_state`, `get_goal_state`, and `navigate_to`. The first two are state queries; `navigate_to` dispatches one deterministic executor step toward the configured goal and returns the applied target pose, progress, moved distance, and stuck-step count. Manipulation exposes `get_gripper_state`, `get_object_state`, `get_target_state`, and `scripted_pick_place_step`. The first three are state queries; `scripted_pick_place_step` advances one phase of a fixed pick-and-place script such as `move_to_pregrasp`, `descend_to_grasp`, or `close_gripper_and_attach`, and returns the applied phase plus the updated observation and motion deltas. The executor therefore exposes task-scoped, deterministic action primitives rather than raw controls or free-form code execution.

D. Metrics and Evaluation Protocol

Success rate is the primary task outcome. An invalid action is an emitted action or tool call that violates the active contract and therefore cannot be dispatched as issued — concretely, a non-dispatchable schema violation such as an unknown tool name or missing required argument. This is distinct from a semantically wrong but schema-valid action (e.g., a valid `navigate_to` call with an incorrect target pose), which passes dispatch but produces an incorrect outcome, and from a task failure, which denotes any run that does not reach the task-complete state regardless of cause. In the structured contract conditions (P1, P2), no semantically-wrong-but-schema-valid actions were observed in the reported matrix; all runs in those conditions reached task-complete success. A recovery is a run that contains an invalid action but still ends successfully after runtime handling. Retries capture the extra repair work used in R1. Planner calls and tool calls together define planner/tool call counts, which is the paper's efficiency proxy.

The planner backend is deterministic by design, so the empirical unit is the task instance under a specified condition rather than repeated stochastic rollouts. The four main comparison cohorts cover 21 fixed task instances (11 navigation easy, 3 manipulation easy, 4 navigation harder, and 3 manipulation harder), evaluated once under all six contract/runtime cells for 126 main-matrix runs. Two focused ablation cohorts add four task instances each to isolate the action-interface and runtime effects, giving 29 task instances and 146 total condition evaluations in the full fixed evaluation set. We therefore report descriptive comparisons over this fixed evaluation set rather than confidence intervals or significance tests across i.i.d. stochastic replications.

The paper draws its quantitative comparisons from the main factorial results plus targeted runtime-only, action-interface ablation, and harder-task summaries. The combined summary provides the aggregate view, while the slice-specific summaries support the ablations and the failure-case interpretation. Within each slice, every compared condition is evaluated on the same fixed tasks, which is why this task coverage still supports a controlled systems comparison even without stochastic rollout replication.

IV. RESULTS

We report the main cohort comparison first, then the mechanism-focused ablations, and finally the harder-task check.

A. Main Contract x Runtime Effect

The main effect is a stable qualitative ordering across navigation easy, tabletop manipulation easy, harder navigation, and harder manipulation. In each slice, P0/R0 is the weakest cell: it fails on all tasks (0/11 navigation easy, 0/3 manipulation easy, 0/4 harder navigation, 0/3 harder manipulation), with every run producing an invalid action. Pairing the same under-specified contract with R1 changes the outcome: P0/R1 succeeds on all tasks (11/11, 3/3, 4/4, 3/3) through recovery after one retry per run. P1/R0, P1/R1, P2/R0, and P2/R1 all succeed on every task with zero invalid actions.

	P0		P1		P2	
Cohort	R0	R1	R0	R1	R0	R1
Navigation Easy	**0/11**	**11/11**	**11/11**	**11/11**	**11/11**	**11/11**
11 runs/cell	× fail	↺ recovered	clean	clean	clean	clean
Manipulation Easy	**0/3**	**3/3**	**3/3**	**3/3**	**3/3**	**3/3**
3 runs/cell	× fail	↺ recovered	clean	clean	clean	clean
Navigation Harder	**0/4**	**4/4**	**4/4**	**4/4**	**4/4**	**4/4**
4 runs/cell	× fail	↺ recovered	clean	clean	clean	clean
Manipulation Harder	**0/3**	**3/3**	**3/3**	**3/3**	**3/3**	**3/3**
3 runs/cell	× fail	↺ recovered	clean	clean	clean	clean

Cell shading encodes status at three grayscale levels: dark = fail (×), medium = recovered (↺), light = clean, so the figure remains readable without color. The harder rows preserve the same ordering as the easy rows; workload growth is summarized separately in Fig. 3 and Table IV.

Fig. 2. Main outcome ordering across the four main cohorts. The under-specified direct-action cell `P0/R0` is consistently weakest, `P0/R1` recovers, and the typed-contract cells remain successful in the evaluated cohorts.

TABLE III
OUTCOME SUMMARY ACROSS THE FOUR MAIN COHORTS. EACH ENTRY KEEPS SUCCESS, INVALID/RUN, AND RETRY/RUN TOGETHER SO THE TABLE STILL DISTINGUISHES RECOVERED `P0/R1` CELLS FROM CLEAN STRUCTURED CELLS.

Cell	**Navigation Easy**	**Manipulation Easy**	**Navigation Harder**	**Manipulation Harder**
P0/R0	0/11; 1.0; 0.0 [fail]	0/3; 1.0; 0.0 [fail]	0/4; 1.0; 0.0 [fail]	0/3; 1.0; 0.0 [fail]
P0/R1	11/11; 1.0; 1.0 [recov.]	3/3; 1.0; 1.0 [recov.]	4/4; 1.0; 1.0 [recov.]	3/3; 1.0; 1.0 [recov.]
P1/R0	11/11; 0.0; 0.0 [clean]	3/3; 0.0; 0.0 [clean]	4/4; 0.0; 0.0 [clean]	3/3; 0.0; 0.0 [clean]
P1/R1	11/11; 0.0; 0.0 [clean]	3/3; 0.0; 0.0 [clean]	4/4; 0.0; 0.0 [clean]	3/3; 0.0; 0.0 [clean]
P2/R0	11/11; 0.0; 0.0 [clean]	3/3; 0.0; 0.0 [clean]	4/4; 0.0; 0.0 [clean]	3/3; 0.0; 0.0 [clean]
P2/R1	11/11; 0.0; 0.0 [clean]	3/3; 0.0; 0.0 [clean]	4/4; 0.0; 0.0 [clean]	3/3; 0.0; 0.0 [clean]

Cell entries are *success; invalid/run; retry/run* followed by overall status: [fail] = no runs succeeded, [recov.] = all succeeded via runtime recovery, [clean] = all succeeded without recovery.

Table III records the same four evaluation cohorts in a compact layout that keeps success, invalid-action exposure, and retries readable without forcing planner/tool counts into every cell. Figure 2 makes the same distinction visually: the recovered `P0/R1` cells are separated from the clean `P1/P2` cells, and the cell text carries the status label so that the comparison does not rely on color alone.

The same comparison separates the roles of contract design and runtime validation. Invalid-action failures are concentrated in the P0 cells, showing that loose task-level verbs are brittle at the planner/executor boundary. Once the planner is forced into a typed tool namespace, the dominant failure mode disappears in the evaluated fixed-R0 ablation. Within the successful structured conditions, P2 matches P1 on success while showing fewer planner/tool calls across the evaluated `P1` versus `P2` comparisons. Table IV isolates that workload comparison so the manuscript can report it without overloading the outcome table. Figure 3 compresses those comparisons into one workload mark per row because planner and tool counts are equal within the plotted `P1/P2` pairs. We interpret that pattern as a workload association within already-strong contracts, not as evidence of a separate success mechanism. Note that planner/tool call counts are a discrete workload proxy, not a latency, token-cost, or wall-clock measurement.

B. Action-Interface and Runtime-Only Ablations

The targeted ablations use different easy-task subsets: the runtime-only *repair-probe cohort* deliberately includes recoverable invalid-first-action probes under fixed P1, whereas the action-interface ablation uses clean easy tasks under fixed R0. The repair-probe cohort was pre-specified from the easy-task pool using a single inclusion criterion: tasks known to elicit a recoverable invalid-first-action emission under P1/R0. This yielded four tasks in total: two navigation and two tabletop-manipulation tasks. Together they isolate the two design axes within the same execution-design framing.

With the contract fixed at P1, moving from R0 to R1 raises success from 2/4 to 4/4 across two navigation tasks and two tabletop manipulation tasks. This slice is intentionally constructed to stress-test recoverable invalid-first-action cases; the 2/4 to 4/4 result should be interpreted as evidence that R1 can repair this specific failure class, not as a general runtime benefit estimate. The mechanism is recovery, not suppression: invalid-action runs remain 2 under both runtimes, but successful invalid-action recoveries rise from 0 to 2 once validation and one retry are enabled. The same direction appears in each family separately, with navigation and manipulation each moving from 0.5 to 1.0 success in this slice. In the runtime traces, the first retry for the navigation probe switches from `move_to_goal` to `get_robot_state`, and the first retry for the manipulation probe switches from `move_object` to `get_object_state`; the repair signal is the explicit error

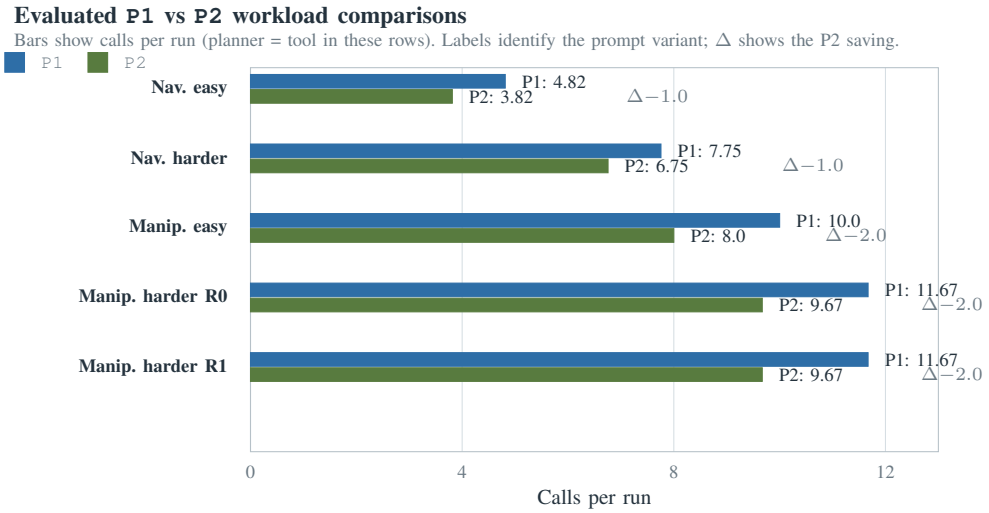

Fig. 3. Workload comparison for the evaluated `P1` versus `P2` rows. Navigation easy and manipulation easy contribute one row each because the `P1`/`P2` workload is identical under `R0` and `R1` there; the harder manipulation slice contributes separate `R0` and `R1` rows. The comparison is descriptive: in this evaluated set, `P2` matches the success profile of `P1` while using fewer planner/tool calls.

TABLE IV
PLANNER/TOOL WORKLOAD ACROSS THE FOUR MAIN COHORTS. FIGURE 3 COMPRESSES THE DUPLICATED PLOTTED PAIRS FOR READABILITY; THIS TABLE PRESERVES THE FULL PER-CELL BREAKDOWN.

Cell	Navigation Easy	Manipulation Easy	Navigation Harder	Manipulation Harder	Δ vs P1 (tool calls)
P0/R0†	1.0/0.0	1.0/0.0	1.0/0.0	1.0/0.0	—
P0/R1	3.82/2.82	9.0/8.0	6.75/5.75	10.67/9.67	—
P1/R0	4.82/4.82	10.0/10.0	7.75/7.75	11.67/11.67	0
P1/R1	4.82/4.82	10.0/10.0	7.75/7.75	11.67/11.67	0
P2/R0	3.82/3.82	8.0/8.0	6.75/6.75	9.67/9.67	−1.0 to −2.0
P2/R1	3.82/3.82	8.0/8.0	6.75/6.75	9.67/9.67	−1.0 to −2.0

Cell entries are *planner calls/tool calls per run*. The Δ column shows the per-cohort range of P2 savings relative to P1.
† P0/R0 reports 0.0 tool calls because all runs fail at the planner stage (invalid action); no tool is dispatched.

string rather than a new second-pass tool schema.

With runtime fixed at R0, P0/R0 fails on all 4/4 evaluated tasks and accumulates four total invalid actions, whereas P1/R0 and P2/R0 succeed on all 4/4 tasks and reduce invalid actions to zero. This establishes an independent contract effect: once the planner must emit dispatchable tool calls, the invalid-action failures that dominate the under-specified interface disappear in this slice. The same comparison also shows the narrower P2 effect: P1/R0 and P2/R0 have the same success profile, but average planner and tool calls each fall from roughly 7.5 to 6.0.

Figure 4 makes both mechanisms visible in one place: the fixed-R0 panel shows typed contracts eliminating invalid actions, while the fixed-P1 runtime panel shows R1 converting the invalid half of the runtime-only slice from failure into recovery, with retries rising from 0.0 to 0.5 per run.

The failure traces identify the dominant P0 failure mode. In navigation, failing P0 runs emit `move_to_goal`; in tabletop manipulation, they emit `move_object`. Both are semantically plausible task verbs, but the executor only dispatches `navigate_to` and `scripted_pick_place_step`. As described in Section III-C, these are injected boundary probes rather than stochastic foundation-model errors. R1 helps because the runtime parses the JSON, validates the tool name and required arguments, and returns the literal error string before one retry. The retry therefore does not introduce a richer second-pass schema; it turns a terminating handoff failure into a recoverable one. Figure 1 includes compact navigation and manipulation traces for exactly that non-dispatchable-to-corrected transition.

a) Isaac Sim cross-backend check.: Two navigation instances were also evaluated under the minimal Isaac Sim stage. The same ordering holds: P0/R0 fails, P0/R1 recovers, P1/P2 succeed with zero invalid actions. This provides a limited cross-backend check that the same boundary ordering also appears in the physics-backed slice.

C. Harder-Task Robustness

The harder-task slices increase planner/tool call counts without changing the condition ordering. In harder navigation, the easy-slice ranking is preserved, while successful cells require roughly three additional planner or tool calls per run. In harder tabletop manipulation, P0/R0 fails on all 3/3 tasks, P0/R1 recovers all 3/3 with one retry per task, and P1/P2 remain fully successful under both runtimes.

A. Fixed-`R0` invalid actions per run (n=2 families)

Action-interface slice: the contract prompt is varied (`P0`/`P1`/`P2`) while the runtime is held at `R0` (no recovery), isolating the effect of typed contracts on invalid-action rates.

Family	P0	P1	P2
Navigation	**1.0** invalid	**0.0** none	**0.0** none
Manipulation	**1.0** invalid	**0.0** none	**0.0** none

Typed contracts remove the invalid-action mode in the fixed-`R0` action-interface slice.

B. Fixed-`P1` runtime outcomes (n=2 families)

Runtime-only slice: the contract is held at `P1` while the runtime is varied (`R0`/`R1`), isolating runtime recovery; invalid/run fixed at 0.5 in both families; right column shows retries/run.

Family	Runtime	Outcomes		Retry
Navigation	R0	clean 1/2	fail 1/2	0.0
	R1	clean 1/2	recov. 1/2	0.5
Manipulation	R0	clean 1/2	fail 1/2	0.0
	R1	clean 1/2	recov. 1/2	0.5

clean success recovered success invalid termination

Fig. 4. Mechanism view from two compact ablation slices (n=4 task instances each, drawn from both navigation and manipulation families). **A** (action-interface slice): under fixed `R0`, `P0` produces one invalid action per run in both families, while `P1`/`P2` reduce that to zero. **B** (runtime-only slice): under fixed `P1`, `R1` turns the invalid half of the runtime-probed runs into recovered success in both families, with retries rising from 0.0 to 0.5 per run.

The P1 versus P2 workload relation also survives the harder manipulation setting. P2 remains about two planner calls and two tool calls per run below P1 under both runtimes while maintaining the same success profile. The harder slices therefore look like cost amplification rather than rank reversal.

V. Discussion and Limitations

A. Discussion

The clearest lesson is that planner-to-executor contract design is an architectural design variable, not a presentation detail. Under the under-specified direct-action contract, the planner often produces task-level verbs that are semantically plausible but not dispatchable under the executor's tool namespace. Explicit contracts instead declare the callable interface up front and reduce the need for executor-side interpretation.

Runtime validation plays a complementary role. It does not replace a well-specified contract, but it can intercept a bad first emission, surface a concrete error, and recover the run with minimal additional machinery. That pattern matters for embodied execution stacks that preserve planner-to-executor separation, including middleware-style action interfaces and guarded execution policies.

At a higher level, the studied boundary maps onto execution architectures already familiar in robotics. The contract specifies the form of a dispatchable request, analogous to how ROS 2 actions define typed request fields and execution interfaces [12]–[14]. Runtime validation then decides whether a malformed or mismatched request should be rejected before execution. A similar separation appears in behavior-tree and TAMP systems, where precondition or guard checks determine whether a proposed action can continue to the executor interface [15]–[18], [20].

The P2 pattern suggests a second, narrower lesson. In the evaluated comparisons, a brief planner-side self-check leaves the success profile unchanged while showing fewer planner/tool calls. The bounded interpretation is not that self-checks are universally best, but that contract design and runtime design remain separable design choices once the callable interface is explicit.

B. Limitations

The conclusions are bounded to a controlled simulator study. The majority of the fixed evaluation matrix runs on a lightweight deterministic reference backend rather than inside a live Isaac Sim physics stage; Isaac Sim serves as a limited qualitative cross-backend check. The evidence therefore covers two task families and is limited to the current easy and harder slices rather than a broad execution benchmark.

The action-interface ablation and runtime-only ablation are compact mechanism checks rather than broad sweeps, and backend coverage differs across slices. The runtime-only ablation is especially specific: it is designed around recoverable invalid-first-action probes, so its strongest claim is about lightweight repair at that boundary.

Using a deterministic mock planner means invalid-action patterns are pre-specified boundary probes rather than stochastic errors from a real foundation model. The empirical frame is therefore descriptive rather than inferential; the deterministic backend and fixed task set preclude confidence intervals or significance tests.

These limits still allow a concrete systems conclusion. In the tested setting, planner-to-executor contracts and lightweight runtime safeguards materially change invalid actions, recovery, and planner/tool call counts. A natural next step is to replicate this factorial with stochastic LLM planner backends and wider scene coverage to test whether the observed ordering generalizes beyond injected boundary probes.

VI. Conclusion

This paper presents a controlled simulator boundary study of how planner-to-executor contract design and runtime validation shape embodied-agent execution, using a lightweight deterministic reference backend for the main factorial and a minimal Isaac Sim navigation slice for qualitative cross-backend checking. In the fixed task matrix studied here, under-specified direct-action handoffs fail at the planner/executor boundary, typed tool calls remove that failure mode in the fixed-R0 slice, and validate-and-retry recovers the recoverable invalid-first-action cases.

The failure analysis shows that the dominant errors are contract violations, not generic task collapse. Explicit contracts prevent those failures at the source, runtime validation repairs

some that remain, and a brief self-check can reduce planner/tool call counts within an already strong contract. These are descriptive comparisons over fixed task instances rather than stochastic rollout estimates. Within the tested setting, the practical implication is to make the callable interface explicit, validate it before dispatch, and treat planner-side self-checks as an efficiency-oriented refinement rather than a separate success mechanism.

REFERENCES

[1] D. Driess, F. Xia, M. S. M. Sajjadi, C. Lynch, A. Chowdhery, B. Ichter, A. Wahid, J. Tompson, Q. Vuong, T. Yu, W. Huang, Y. Chebotar, P. Sermanet, D. Duckworth, S. Levine, V. Vanhoucke, K. Hausman, M. Toussaint, K. Greff, A. Zeng, I. Mordatch, and P. Florence, "PaLM-E: An embodied multimodal language model," in *Proceedings of the 40th International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, vol. 202. PMLR, 2023, pp. 8469–8488. [Online]. Available: https://proceedings.mlr.press/v202/driess23a.html

[2] B. Zitkovich, T. Yu, S. Xu, P. Xu, T. Xiao, F. Xia, J. Wu, P. Wohlhart, S. Welker, A. Wahid, Q. Vuong, V. Vanhoucke, H. Tran, R. Soricut, A. Singh, J. Singh, P. Sermanet, P. R. Sanketi, G. Salazar, M. S. Ryoo, K. Reymann, K. Rao, K. Pertsch, I. Mordatch, H. Michalewski, Y. Lu, S. Levine, L. Lee, T.-W. E. Lee, I. Leal, Y. Kuang, D. Kalashnikov, R. Julian, N. J. Joshi, A. Irpan, B. Ichter, J. Hsu, A. Herzog, K. Hausman, K. Gopalakrishnan, C. Fu, P. Florence, C. Finn, K. A. Dubey, D. Driess, T. Ding, K. M. Choromanski, X. Chen, Y. Chebotar, J. Carbajal, N. Brown, A. Brohan, M. G. Arenas, and K. Han, "RT-2: Vision-language-action models transfer web knowledge to robotic control," in *Proceedings of The 7th Conference on Robot Learning*, ser. Proceedings of Machine Learning Research, vol. 229. PMLR, 2023, pp. 2165–2183. [Online]. Available: https://proceedings.mlr.press/v229/zitkovich23a.html

[3] R. Yang, H. Chen, J. Zhang, M. Zhao, C. Qian, K. Wang, Q. Wang, T. V. Koripella, M. Movahedi, M. Li, H. Ji, H. Zhang, and T. Zhang, "EmbodiedBench: Comprehensive benchmarking multi-modal large language models for vision-driven embodied agents," in *Proceedings of the 42nd International Conference on Machine Learning*, ser. Proceedings of Machine Learning Research, vol. 267. PMLR, 2025, pp. 70 576–70 631. [Online]. Available: https://proceedings.mlr.press/v267/yang25f.html

[4] X. Lu, Z. Chen, X. Hu, Y. Zhou, W. Zhang, D. Liu, L. Sheng, and J. Shao, "IS-Bench: Evaluating interactive safety of VLM-driven embodied agents in daily household tasks," arXiv preprint arXiv:2506.16402, 2025. [Online]. Available: https://arxiv.org/abs/2506.16402

[5] N. Kachaev, A. Spiridonov, A. Gorodetsky, K. Muravyev, N. Oskolkov, A. Narendra, V. Shakhuro, D. Makarov, A. I. Panov, P. Fedotova, and A. K. Kovalev, "Mind and motion aligned: A joint evaluation IsaacSim benchmark for task planning and low-level policies in mobile manipulation," arXiv preprint arXiv:2508.15663, 2025. [Online]. Available: https://arxiv.org/abs/2508.15663

[6] W. Huang, F. Xia, T. Xiao, H. Chan, J. Liang, P. Florence, A. Zeng, J. Tompson, I. Mordatch, Y. Chebotar, P. Sermanet, N. Brown, T. Jackson, L. Luu, S. Levine, K. Hausman, and B. Ichter, "Inner monologue: Embodied reasoning through planning with language models," in *Proceedings of The 6th Conference on Robot Learning*, ser. Proceedings of Machine Learning Research, vol. 205. PMLR, 2023, pp. 1769–1782. [Online]. Available: https://proceedings.mlr.press/v205/huang23c.html

[7] W. Huang, C. Wang, R. Zhang, Y. Li, J. Wu, and L. Fei-Fei, "VoxPoser: Composable 3D value maps for robotic manipulation with language models," in *Proceedings of The 7th Conference on Robot Learning*, ser. Proceedings of Machine Learning Research, vol. 229. PMLR, 2023, pp. 540–562. [Online]. Available: https://proceedings.mlr.press/v229/huang23b.html

[8] P. Liu, Y. Orru, J. Vakil, C. Paxton, N. M. M. Shafiullah, and L. Pinto, "OK-Robot: What really matters in integrating open-knowledge models for robotics," arXiv preprint arXiv:2401.12202, 2024. [Online]. Available: https://arxiv.org/abs/2401.12202

[9] M. Ahn, A. Brohan, N. Brown, Y. Chebotar, O. Cortes, B. David, C. Finn, C. Fu, K. Gopalakrishnan, K. Hausman, A. Herzog, D. Ho, J. Hsu, J. Ibarz, B. Ichter, A. Irpan, E. Jang, R. J. Ruano, K. Jeffrey, S. Jesmonth, N. J. Joshi, R. Julian, D. Kalashnikov, Y. Kuang, K.-H. Lee, S. Levine, Y. Lu, L. Luu, C. Parada, P. Pastor, J. Quiambao, K. Rao, J. Rettinghouse, D. Reyes, P. Sermanet, N. Sievers, C. Tan, A. Toshev, V. Vanhoucke, F. Xia, T. Xiao, P. Xu, S. Xu, M. Yan, and A. Zeng, "Do as I can, not as I say: Grounding language in robotic affordances," in *Proceedings of The 6th Conference on Robot Learning*, ser. Proceedings of Machine Learning Research, vol. 205. PMLR, 2023, pp. 287–318. [Online]. Available: https://proceedings.mlr.press/v205/ichter23a.html

[10] J. Liang, W. Huang, F. Xia, P. Xu, K. Hausman, B. Ichter, P. Florence, and A. Zeng, "Code as policies: Language model programs for embodied control," in *2023 IEEE International Conference on Robotics and Automation (ICRA)*, 2023, pp. 9493–9500. [Online]. Available: https://doi.org/10.1109/ICRA48891.2023.10160591

[11] I. Singh, V. Blukis, A. Mousavian, A. Goyal, D. Xu, J. Tremblay, D. Fox, J. Thomason, and A. Garg, "Progprompt: Generating situated robot task plans using large language models," *Autonomous Robots*, vol. 47, pp. 999–1012, 2023. [Online]. Available: https://doi.org/10.1007/s10514-023-10135-3

[12] S. Macenski, T. Foote, B. Gerkey, C. Lalancette, and W. Woodall, "Robot operating system 2: Design, architecture, and uses in the wild," *Science Robotics*, vol. 7, no. 66, 2022. [Online]. Available: https://doi.org/10.1126/scirobotics.abm6074

[13] G. Biggs, J. Perron, and S. Loretz, "ROS 2 actions," ROS 2 design article, 2019, last modified 2020-05. [Online]. Available: https://design.ros2.org/articles/actions.html

[14] F. Martin, J. G. Clavero, V. Matellan, and F. J. Rodriguez, "PlanSys2: A planning system framework for ROS2," in *2021 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2021, pp. 9742–9749. [Online]. Available: https://doi.org/10.1109/IROS51168.2021.9636544

[15] M. Colledanchise and P. Ögren, *Behavior Trees in Robotics and AI*. CRC Press, 2018. [Online]. Available: https://doi.org/10.1201/9780429489105

[16] BehaviorTree.CPP Documentation, "BehaviorTree.CPP docs," Official documentation, version 4.8; accessed 2026-03-18. [Online]. Available: https://www.behaviortree.dev/docs/intro/

[17] C. R. Garrett, T. Lozano-Pérez, and L. P. Kaelbling, "PDDLStream: Integrating symbolic planners and blackbox samplers via optimistic adaptive planning," *Proceedings of the International Conference on Automated Planning and Scheduling*, vol. 30, pp. 440–448, 2020. [Online]. Available: https://doi.org/10.1609/ICAPS.V30I1.6739

[18] C. R. Garrett, R. Chitnis, R. Holladay, B. Kim, T. Silver, L. P. Kaelbling, and T. Lozano-Pérez, "Integrated task and motion planning," *Annual Review of Control, Robotics, and Autonomous Systems*, vol. 4, no. 1, pp. 265–293, 2021. [Online]. Available: https://doi.org/10.1146/annurev-control-091420-084139

[19] NVIDIA, "Isaac Sim documentation," Official documentation, accessed 2026-03-17. [Online]. Available: https://docs.isaacsim.omniverse.nvidia.com/latest/index.html

[20] BehaviorTree.CPP Documentation, "Pre and post conditions," Official documentation, version 4.0.2; accessed 2026-03-18. [Online]. Available: https://www.behaviortree.dev/docs/4.0.2/tutorial-advanced/pre_post_conditions/